Date	Name	Domain	Focus	Keywords	Task Types	Metrics	Models	Citation
2020-09-07	MMLU (Massive Multitask Language Understanding)	Multidomain	Academic knowledge and reasoning across 57 subjects	multitask, multiple-choice, zero-shot, few-shot, knowledge probing	Multiple choice	Accuracy	GPT-4o, Gemini 1.5 Pro, o1, DeepSeek-R1	**hendrycks2021measuring**⇒
2023-11-20	GPQA Diamond	Science	Graduate-level scientific reasoning	Google-proof, graduate-level, science QA, chemistry, physics	Multiple choice, Multi-step QA	Accuracy	o1, DeepSeek-R1	[1]⇒
2018-03-14	ARC-Challenge (Advanced Reasoning Challenge)	Science	Grade-school science with reasoning emphasis	grade-school, science QA, challenge set, reasoning	Multiple choice	Accuracy	GPT-4, Claude	[2]⇒
2025-01-24	Humanity's Last Exam	Multidomain	Broad cross-domain academic reasoning	cross-domain, academic exam, multiple-choice, multi-disciplinary	Multiple choice	Accuracy		[3]⇒
2024-11-07	FrontierMath	Mathematics	Challenging advanced mathematical reasoning	symbolic reasoning, number theory, algebraic geometry, category theory	Problem solving	Accuracy		[4]⇒
2024-07-18	SciCode	Scientific Programming	Scientific code generation and problem solving	code synthesis, scientific computing, programming benchmark	Coding	Solve rate (%)	Claude3.5-Sonnet	[5]⇒
2025-03-13	AIME (American Invitational Mathematics Examination)	Mathematics	Pre-college advanced problem solving	algebra, combinatorics, number theory, geometry	Problem solving	Accuracy		[6]⇒
2025-02-15	MATH-500	Mathematics	Math reasoning generalization	calculus, algebra, number theory, geometry	Problem solving	Accuracy		[7]⇒

Continued on next page

Date	Name	Domain	Focus	Keywords	Task Types	Metrics	Models	Citation
2024-04-02	CURIE (Scientific Long-Context Understanding, Reasoning and Information Extraction)	Multidomain Science	Long-context scientific reasoning	long-context, information extraction, multimodal	Information extraction, Reasoning, Concept tracking, Aggregation, Algebraic manipulation, Multimodal comprehension	Accuracy		[8]⇒
2023-01-26	FEABench (Finite Element Analysis Benchmark)	Computational Engineering	FEA simulation accuracy and performance	finite element, simulation, PDE	Simulation, Performance evaluation	Solve time, Error norm	FEniCS, deal.II	[9]⇒
2024-07-12	SPIQA (Scientific Paper Image Question Answering)	Computer Science	Multimodal QA on scientific figures	multimodal QA, figure understanding, table comprehension, chain-of-thought	Question answering, Multimodal QA, Chain-of-Thought evaluation	Accuracy, F1 score	Chain-of-Thought models, Multimodal QA systems	[10]⇒
2020-09-28	MedQA	Medical Question Answering	Medical board exam QA	USMLE, diagnostic QA, medical knowledge, multilingual	Multiple choice	Accuracy	Neural reader, Retrieval-based QA systems	[11]⇒
2025-05-13	BaisBench (Biological AI Scientist Benchmark)	Computational Biology	Omics-driven AI research tasks	single-cell annotation, biological QA, autonomous discovery	Cell type annotation, Multiple choice	Annotation accuracy, QA accuracy	LLM-based AI scientist agents	[12]⇒
2023-01-26	MOLGEN	Computational Chemistry	Molecular generation and optimization	SELFIES, GAN, property optimization	Distribution learning, Goal-oriented generation	Validity%, Novelty%, QED, Docking score	MolGen	[13]⇒
2020-05-02	Open Graph Benchmark (OGB) - Biology	Graph ML	Biological graph property prediction	node prediction, link prediction, graph classification	Node property prediction, Link property prediction, Graph property prediction	Accuracy, ROC-AUC	GCN, GraphSAGE, GAT	[14]⇒
2011-10-01	Materials Project	Materials Science	DFT-based property prediction	DFT, materials genome, high-throughput	Property prediction	MAE, R^2	Automatminer, Crystal Graph Neural Networks	[15]⇒

Continued on next page

Date	Name	Domain	Focus	Keywords	Task Types	Metrics	Models	Citation
2020-10-20	OCP (Open Catalyst Project)	Chemistry; Materials Science	Catalyst adsorption energy prediction	DFT relaxations, adsorption energy, graph neural networks	Energy prediction, Force prediction	MAE (energy), MAE (force)	CGCNN, SchNet, DimeNet++, GemNet-OC	[16]–[19]⇒
2023-06-20	JARVIS-Leaderboard	Materials Science; Benchmarking	Comparative evaluation of materials design methods	leaderboards, materials methods, simulation	Method benchmarking, Leaderboard ranking	MAE, RMSE, Accuracy		[20]⇒
2022-02-22	Quantum Computing Benchmarks (QML)	Quantum Computing	Quantum algorithm performance evaluation	quantum circuits, state preparation, error correction	Circuit benchmarking, State classification	Fidelity, Success probability	IBM Q, IonQ, AQT@LBNL	[21]⇒
2024-10-01	CFDBench (Fluid Dynamics)	Fluid Dynamics; Scientific ML	Neural operator surrogate modeling	neural operators, CFD, FNO, DeepONet	Surrogate modeling	L2 error, MAE	FNO, DeepONet, U-Net	[22]⇒
	SatImgNet	Remote Sensing	Satellite imagery classification	land-use, zero-shot, multi-task	Image classification	Accuracy		[23]⇒
2023-07-19	ClimateLearn	Climate Science; Forecasting	ML for weather and climate modeling	medium-range forecasting, ERA5, data-driven	Forecasting	RMSE, Anomaly correlation	CNN baselines, ResNet variants	[24]⇒
2022-06-09	BIG-Bench (Beyond the Imitation Game Benchmark)	NLP; AI Evaluation	Diverse reasoning and generalization tasks	few-shot, multi-task, bias analysis	Few-shot evaluation, Multi-task evaluation	Accuracy, Task-specific metrics	GPT-3, Dense Transformers, Sparse Transformers	[25]⇒
2019-11-20	CommonSenseQA	NLP; Commonsense	Commonsense question answering	ConceptNet, multiple-choice, adversarial	Multiple choice	Accuracy	BERT-large, RoBERTa, GPT-3	[26]⇒
2019-07-24	Winogrande	NLP; Commonsense	Winograd Schema-style pronoun resolution	adversarial, pronoun resolution	Pronoun resolution	Accuracy, AUC	RoBERTa, BERT, GPT-2	[27]⇒
2024-05-01	Jet Classification	Particle Physics	Real-time classification of particle jets using HL-LHC simulation features	classification, real-time ML, jet tagging, QKeras	Classification	Accuracy, AUC	Keras DNN, QKeras quantized DNN	[28]⇒

Continued on next page

Date	Name	Domain	Focus	Keywords	Task Types	Metrics	Models	Citation
2024-05-01	Irregular Sensor Data Compression	Particle Physics	Real-time compression of sparse sensor data with autoencoders	compression, autoencoder, sparse data, irregular sampling	Compression	MSE, Compression ratio	Autoencoder, Quantized autoencoder	[29]⇒
2024-05-01	Beam Control	Accelerators and Magnets	Reinforcement learning control of accelerator beam position	RL, beam stabilization, control systems, simulation	Control	Stability, Control loss	DDPG, PPO (planned)	[30], [31]⇒
2024-07-08	Ultrafast jet classification at the HL-LHC	Particle Physics	FPGA-optimized real-time jet origin classification at the HL-LHC	jet classification, FPGA, quantization-aware training, Deep Sets, Interaction Networks	Classification	Accuracy, Latency, Resource utilization	MLP, Deep Sets, Interaction Network	
2024-10-15	Quench detection	Accelerators and Magnets	Real-time detection of superconducting magnet quenches using ML	quench detection, autoencoder, anomaly detection, real-time	Anomaly detection, Quench localization	ROC-AUC, Detection latency	Autoencoder, RL agents (in development)	
2024-10-15	DUNE	Particle Physics	Real-time ML for DUNE DAQ time-series data	DUNE, time-series, real-time, trigger	Trigger selection, Time-series anomaly detection	Detection efficiency, Latency	CNN, LSTM (planned)	
2025-01-08	Intelligent experiments through real-time AI	Instrumentation and Detectors; Nuclear Physics; Particle Physics	Real-time FPGA-based triggering and detector control for sPHENIX and future EIC	FPGA, Graph Neural Network, hls4ml, real-time inference, detector control	Trigger classification, Detector control, Real-time inference	Accuracy (charm and beauty detection), Latency (µs), Resource utilization (LUT/FF/BRAM/DSP)	Bipartite Graph Network with Set Transformers (BGN-ST), GarNet (edge-classifier)	[32]⇒
2025-01-09	Neural Architecture Codesign for Fast Physics Applications	Physics; Materials Science; Particle Physics	Automated neural architecture search and hardware-efficient model codesign for fast physics applications	neural architecture search, FPGA deployment, quantization, pruning, hls4ml	Classification, Peak finding	Accuracy, Latency, Resource utilization	NAC-based BraggNN, NAC-optimized Deep Sets (jet)	[33]⇒
2024-06-24	Smart Pixels for LHC	Particle Physics; Instrumentation and Detectors	On-sensor, in-pixel ML filtering for high-rate LHC pixel detectors	smart pixel, on-sensor inference, data reduction, trigger	Image Classification, Data filtering	Data rejection rate, Power per pixel	2-layer pixel NN	[34]⇒

Continued on next page

Date	Name	Domain	Focus	Keywords	Task Types	Metrics	Models	Citation
2023-10-03	HEDM (BraggNN)	Material Science	Fast Bragg peak analysis using deep learning in diffraction microscopy	BraggNN, diffraction, peak finding, HEDM	Peak detection	Localization accuracy, Inference time	BraggNN	[35]⇒
2023-12-03	4D-STEM	Material Science	Real-time ML for scanning transmission electron microscopy	4D-STEM, electron microscopy, real-time, image processing	Image Classification, Streamed data inference	Classification accuracy, Throughput	CNN models (prototype)	[36]⇒
2023-12-05	In-Situ High-Speed Computer Vision	Fusion/Plasma	Real-time image classification for in-situ plasma diagnostics	plasma, in-situ vision, real-time ML	Image Classification	Accuracy, FPS	CNN	[37]⇒
2020-01-01	BenchCouncil AIBench	General	End-to-end AI benchmarking across micro, component, and application levels	benchmarking, AI systems, application-level evaluation	Training, Inference, End-to-end AI workloads	Throughput, Latency, Accuracy	ResNet, BERT, GANs, Recommendation systems	[38]⇒
2020-01-01	BenchCouncil BigDataBench	General	Big data and AI benchmarking across structured, semi-structured, and unstructured data workloads	big data, AI benchmarking, data analytics	Data preprocessing, Inference, End-to-end data pipelines	Data throughput, Latency, Accuracy	CNN, LSTM, SVM, XGBoost	[39]⇒
2021-10-20	MLPerf HPC	Cosmology, Climate, Protein Structure, Catalysis	Scientific ML training and inference on HPC systems	HPC, training, inference, scientific ML	Training, Inference	Training time, Accuracy, GPU utilization	CosmoFlow, DeepCAM, OpenCatalyst	[40]⇒
2023-06-01	MLCommons Science	Earthquake, Satellite Image, Drug Discovery, Electron Microscope, CFD	AI benchmarks for scientific applications including time-series, imaging, and simulation	science AI, benchmark, MLCommons, HPC	Time-series analysis, Image classification, Simulation surrogate modeling	MAE, Accuracy, Speedup vs simulation	CNN, GNN, Transformer	[41]⇒
2021-07-05	LHC New Physics Dataset	Particle Physics; Real-time Triggering	Real-time LHC event filtering for anomaly detection using proton collision data	anomaly detection, proton collision, real-time inference, event filtering, unsupervised ML	Anomaly detection, Event classification	ROC-AUC, Detection efficiency	Autoencoder, Variational autoencoder, Isolation forest	[42]⇒

Continued on next page

Date	Name	Domain	Focus	Keywords	Task Types	Metrics	Models	Citation
2023-07-17	MLCommons Medical AI	Healthcare; Medical AI	Federated benchmarking and evaluation of medical AI models across diverse real-world clinical data	medical AI, federated evaluation, privacy-preserving, fairness, healthcare benchmarks	Federated evaluation, Model validation	ROC AUC, Accuracy, Fairness metrics	MedPerf-validated CNNs, GaN-DLF workflows	**karargyris2023federated**⇒
2024-10-28	CaloChallenge 2022	LHC Calorimeter; Particle Physics	Fast generative-model-based calorimeter shower simulation evaluation	calorimeter simulation, generative models, surrogate modeling, LHC, fast simulation	Surrogate modeling	Histogram similarity, Classifier AUC, Generation latency	VAE variants, GAN variants, Normalizing flows, Diffusion models	[43]⇒
ongoing	Papers With Code (SOTA Platform)	General ML; All domains	Open platform tracking state-of-the-art results, benchmarks, and implementations across ML tasks and papers	leaderboard, benchmarking, reproducibility, open-source	Multiple (Classification, Detection, NLP, etc.)	Task-specific (Accuracy, F1, BLEU, etc.)	All published models with code	[44]⇒
2022-01-01	Codabench	General ML; Multiple	Open-source platform for organizing reproducible AI benchmarks and competitions	benchmark platform, code submission, competitions, meta-benchmark	Multiple	Submission count, Leaderboard ranking, Task-specific metrics	Arbitrary code submissions	[45]⇒
2021-09-27	Sabath (SBI-FAIR)	Systems; Metadata	FAIR metadata framework for ML-driven surrogate workflows in HPC systems	meta-benchmark, metadata, HPC, surrogate modeling	Systems benchmarking	Metadata completeness, FAIR compliance	N/A	[46]⇒
2022-10-13	PDEBench	CFD; Weather Modeling	Benchmark suite for ML-based surrogates solving time-dependent PDEs	PDEs, CFD, scientific ML, surrogate modeling, NeurIPS	Supervised Learning	RMSE, boundary RMSE, Fourier RMSE	FNO, U-Net, PINN, Gradient-Based inverse methods	[47]⇒

Continued on next page

Date	Name	Domain	Focus	Keywords	Task Types	Metrics	Models	Citation
2024-12-03	The Well	biological systems, fluid dynamics, acoustic scattering, astrophysical MHD	Foundation model + surrogate dataset spanning 16 physical simulation domains	surrogate modeling, foundation model, physics simulations, spatiotemporal dynamics	Supervised Learning	Dataset size, Domain breadth	FNO baselines, U-Net baselines	**ohana2024well**⇒
2024-10-31	LLM-Inference-Bench	LLM; HPC/inference	Hardware performance benchmarking of LLMs on AI accelerators	LLM, inference benchmarking, GPU, accelerator, throughput	Inference Benchmarking	Token throughput (tok/s), Latency, Framework-hardware mix performance	LLaMA-2-7B, LLaMA-2-70B, Mistral-7B, Qwen-7B	**chitty2024llm**⇒
2023-12-12	SGLang Framework	LLM Vision	Fast serving framework for LLMs and vision-language models	LLM serving, vision-language, RadixAttention, performance, JSON decoding	Model serving framework	Tokens/sec, Time-to-first-token, Throughput gain vs baseline	LLaVA, DeepSeek, Llama	[48]⇒
2023-09-12	vLLM Inference and Serving Engine	LLM; HPC/inference	High-throughput, memory-efficient inference and serving engine for LLMs	LLM inference, PagedAttention, CUDA graph, streaming API, quantization	Inference Benchmarking	Tokens/sec, Time to First Token (TTFT), Memory footprint	LLaMA, Mixtral, FlashAttention-based models	[49]⇒
2022-06-22	vLLM Performance Dashboard	LLM; HPC/inference	Interactive dashboard showing inference performance of vLLM	Dashboard, Throughput visualization, Latency analysis, Metric tracking	Performance visualization	Tokens/sec, TTFT, Memory usage	LLaMA-2, Mistral, Qwen	[50]⇒
2022-04-01	Nixtla NeuralForecast	Time-series forecasting; General ML	High-performance neural forecasting library with >30 models	time-series, neural forecasting, NBEATS, NHITS, TFT, probabilistic forecasting, usability	Time-series forecasting	RMSE, MAPE, CRPS	NBEATS, NHITS, TFT, DeepAR	[51]⇒
2023-06-01	Nixtla Neural Forecast NHITS	Time-series; General ML	Official NHITS implementation for long-horizon time series forecasting	NHITS, long-horizon forecasting, neural interpolation, time-series	Time-series forecasting	RMSE, MAPE	NHITS	[52]⇒

Continued on next page

Date	Name	Domain	Focus	Keywords	Task Types	Metrics	Models	Citation
2023-10-03	Nixtla Neural Forecast TimeLLM	Time-series; General ML	Reprogramming LLMs for time series forecasting	Time-LLM, language model, time-series, reprogramming	Time-series forecasting	RMSE, MAPE	Time-LLM	**jin2023time⇒**
2023-10-05	Nixtla Neural Forecast TimeGPT	Time-series; General ML	Time-series foundation model "TimeGPT" for forecasting and anomaly detection	TimeGPT, foundation model, time-series, generative model	Time-series forecasting, Anomaly detection	RMSE, Anomaly detection metrics	TimeGPT	[53]⇒
2025-03-03	HDR ML Anomaly Challenge (Gravitational Waves)	Astrophysics; Time-series	Detecting anomalous gravitational-wave signals from LIGO/Virgo datasets	anomaly detection, gravitational waves, astrophysics, time-series	Anomaly detection	ROC-AUC, Precision/Recall	Deep latent CNNs, Autoencoders	[54]⇒
2025-03-03	HDR ML Anomaly Challenge (Butterfly)	Genomics; Image/CV	Detecting hybrid butterflies via image anomaly detection in genomic-informed dataset	anomaly detection, computer vision, genomics, butterfly hybrids	Anomaly detection	Classification accuracy, F1 score	CNN-based detectors	[55]⇒
2025-03-03	HDR ML Anomaly Challenge (Sea Level Rise)	Climate Science; Time-series, Image/CV	Detecting anomalous sea-level rise and flooding events via time-series and satellite imagery	anomaly detection, climate science, sea-level rise, time-series, remote sensing	Anomaly detection	ROC-AUC, Precision/Recall	CNNs, RNNs, Transformers	[56]⇒
2025-01-24	Single Qubit Readout on QICK System	Quantum Computing	Real-time single-qubit state classification using FPGA firmware	qubit readout, hls4ml, FPGA, QICK	Classification	Accuracy, Latency	hls4ml quantized NN	[57]⇒
2023-11-20	GPQA: A Graduate-Level Google-Proof Question and Answer Benchmark	Science (Biology, Physics, Chemistry)	Graduate-level, expert-validated multiple-choice questions hard even with web access	Google-proof, multiple-choice, expert reasoning, science QA	Multiple choice	Accuracy	GPT-4 baseline	[58]⇒

Continued on next page

Date	Name	Domain	Focus	Keywords	Task Types	Metrics	Models	Citation
2024-12-13	SeafloorAI	Marine Science; Vision-Language	Large-scale vision-language dataset for seafloor mapping and geological classification	sonar imagery, vision-language, seafloor mapping, segmentation, QA	Image segmentation, Vision-language QA	Segmentation pixel accuracy, QA accuracy	SegFormer, ViLT-style multimodal models	[59]⇒
2024-12-13	SuperCon3D	Materials Science; Superconductivity	Dataset and models for predicting and generating high-Tc superconductors using 3D crystal structures	superconductivity, crystal structures, equivariant GNN, generative models	Regression (Tc prediction), Generative modeling	MAE (Tc), Validity of generated structures	SODNet, DiffCSP-SC	**zhuang2024supercon3d**⇒
2024-12-13	GeSS	Scientific ML; Geometric Deep Learning	Benchmark suite evaluating geometric deep learning models under real-world distribution shifts	geometric deep learning, distribution shift, OOD robustness, scientific applications	Classification, Regression	Accuracy, RMSE, OOD robustness delta	GCN, EGNN, DimeNet++	**zou2024gess**⇒
2024-12-13	Vocal Call Locator (VCL)	Neuroscience; Bioacoustics	Benchmarking sound-source localization of rodent vocalizations from multi-channel audio	source localization, bioacoustics, time-series, SSL	Sound source localization	Localization error (cm), Recall/Precision	CNN-based SSL models	[60]⇒
2024-12-13	MassSpecGym	Cheminformatics; Molecular Discovery	Benchmark suite for discovery and identification of molecules via MS/MS	mass spectrometry, molecular structure, de novo generation, retrieval, dataset	De novo generation, Retrieval, Simulation	Structure accuracy, Retrieval precision, Simulation MSE	Graph-based generative models, Retrieval baselines	[61]⇒
2024-12-13	Urban Data Layer (UDL)	Urban Computing; Data Engineering	Unified data pipeline for multi-modal urban science research	data pipeline, urban science, multi-modal, benchmark	Prediction, Classification	Task-specific accuracy or RMSE	Baseline regression/classification pipelines	[62]⇒

Continued on next page

Date	Name	Domain	Focus	Keywords	Task Types	Metrics	Models	Citation
2024-12-13	Delta Squared-DFT	Computational Chemistry; Materials Science	Benchmarking machine-learning corrections to DFT using Delta Squared-trained models for reaction energies	density functional theory, Delta Squared-ML correction, reaction energetics, quantum chemistry	Regression	Mean Absolute Error (eV), Energy ranking accuracy	Delta Squared-ML correction networks, Kernel ridge regression	[63]⇒
2024-12-13	LLMs for Crop Science	Agricultural Science; NLP	Evaluating LLMs on crop trait QA and textual inference tasks with domain-specific prompts	crop science, prompt engineering, domain adaptation, question answering	Question Answering, Inference	Accuracy, F1 score	GPT-4, LLaMA-2-13B, T5-XXL	[64]⇒
2024-12-13	SPIQA (LLM)	Multimodal Scientific QA; Computer Vision	Evaluating LLMs on image-based scientific paper figure QA tasks (LLM Adapter performance)	multimodal QA, scientific figures, image+text, chain-of-thought prompting	Multimodal QA	Accuracy, F1 score	LLaVA, MiniGPT-4, Owl-LLM adapter variants	[65]⇒

References

[1] D. Rein, B. L. Hou, A. C. Stickland, *et al.*, *Gpqa: A graduate-level google-proof q and a benchmark*, 2023. [Online]. Available: https://arxiv.org/abs/2311.12022.

[2] P. Clark, I. Cowhey, O. Etzioni, *et al.*, "Think you have solved question answering? try arc, the ai2 reasoning challenge," in *EMNLP 2018*, 2018, pp. 237–248. [Online]. Available: https://allenai.org/data/arc.

[3] L. Phan, A. Gatti, Z. Han, *et al.*, *Humanity's last exam*, 2025. [Online]. Available: https://arxiv.org/abs/2501.14249.

[4] E. Glazer, E. Erdil, T. Besiroglu, *et al.*, *Frontiermath: A benchmark for evaluating advanced mathematical reasoning in ai*, 2024. [Online]. Available: https://arxiv.org/abs/2411.04872.

[5] M. Tian, L. Gao, S. Zhang, *et al.*, *Scicode: A research coding benchmark curated by scientists*, 2024. [Online]. Available: https://arxiv.org/abs/2407.13168.

[6] TBD, *Aime*, [Online accessed 2025-06-24], Mar. 2025. [Online]. Available: https://www.vals.ai/benchmarks/aime-2025-03-13.

[7] HuggingFaceH4, *Math-500*, 2025. [Online]. Available: https://huggingface.co/datasets/HuggingFaceH4/MATH-500.

[8] T. A. authors, *Scientific reasoning benchmarks from the curie dataset*, 2024. [Online]. Available: https://arxiv.org/abs/2404.02029.

[9] A. Institute, *Feabench: A finite element analysis benchmark*, 2023. [Online]. Available: https://github.com/alleninstitute/feabench.

[10] X. Zhong, Y. Gao, and S. Gururangan, *Spiqa: Scientific paper image question answering*, 2024. [Online]. Available: https://arxiv.org/abs/2407.09413.

[11] D. Jin, Y. Li, Y. Zhang, *et al.*, *What disease does this patient have? a large-scale open-domain question answering dataset from medical exams*, 2020. [Online]. Available: https://arxiv.org/abs/2009.13081.

[12] E. Luo, J. Jia, Y. Xiong, *et al.*, *Benchmarking ai scientists in omics data-driven biological research*, 2025. [Online]. Available: https://arxiv.org/abs/2505.08341.

[13] Y. Fang, N. Zhang, Z. Chen, *et al.*, *Domain-agnostic molecular generation with chemical feedback*, 2023. [Online]. Available: https://arxiv.org/abs/2301.11259.

[14] W. Hu, M. Fey, M. Zitnik, *et al.*, *Open graph benchmark: Datasets for machine learning on graphs*, 2020. [Online]. Available: https://arxiv.org/abs/2005.00687.

[15] A. Jain, S. P. Ong, G. Hautier, *et al.*, "The materials project: A materials genome approach," *APL Materials*, vol. 1, no. 1, 2013. DOI: 10.1063/1.4812323. [Online]. Available: https://materialsproject.org/.

[16] L. Chanussot, A. Das, S. Goyal, *et al.*, "The open catalyst 2020 (oc20) dataset and community challenges," *ACS Catalysis*, vol. 11, no. 10, pp. 6059–6072, 2021. DOI: 10.1021/acscatal.0c04525. [Online]. Available: https://pubs.acs.org/doi/10.1021/acscatal.0c04525.

[17] R. Tran, J. Lan, M. Shuaibi, *et al.*, "The open catalyst 2022 (oc22) dataset and challenges for oxide electrocatalysts," *ACS Catalysis*, vol. 13, no. 5, pp. 3066–3084, 2023. DOI: 10.1021/acscatal.2c05426. [Online]. Available: https://pubs.acs.org/doi/10.1021/acscatal.2c05426.

[18] L. Chanussot, A. Das, S. Goyal, *et al.*, "Open catalyst 2020 (oc20) dataset and community challenges," *ACS Catalysis*, vol. 11, no. 10, pp. 6059–6072, 2021. DOI: 10.1021/acscatal.0c04525. eprint: https://doi.org/10.1021/acscatal.0c04525. [Online]. Available: https://doi.org/10.1021/acscatal.0c04525.

[19] R. Tran, J. Lan, M. Shuaibi, *et al.*, "The open catalyst 2022 (oc22) dataset and challenges for oxide electrocatalysts," *ACS Catalysis*, vol. 13, no. 5, pp. 3066–3084, Feb. 2023, ISSN: 2155-5435. DOI: 10.1021/acscatal.2c05426. [Online]. Available: http://dx.doi.org/10.1021/acscatal.2c05426.

[20] K. Choudhary, D. Wines, K. Li, *et al.*, "JARVIS-Leaderboard: A large scale benchmark of materials design methods," *npj Computational Materials*, vol. 10, no. 1, p. 93, 2024. DOI: 10.1038/s41524-024-01259-w. [Online]. Available: https://doi.org/10.1038/s41524-024-01259-w.

[21] F. J. Kiwit, M. Marso, P. Ross, C. A. Riofrío, J. Klepsch, and A. Luckow, "Application-oriented benchmarking of quantum generative learning using quark," in *2023 IEEE International Conference on Quantum Computing and Engineering (QCE)*, IEEE, Sep. 2023, pp. 475–484. DOI: 10.1109/qce57702.2023.00061. [Online]. Available: http://dx.doi.org/10.1109/QCE57702.2023.00061.

[22] Y. Luo, Y. Chen, and Z. Zhang, *Cfdbench: A large-scale benchmark for machine learning methods in fluid dynamics*, 2024. [Online]. Available: https://arxiv.org/abs/2310.05963.

[23] J. Roberts, K. Han, and S. Albanie, *Satin: A multi-task metadataset for classifying satellite imagery using vision-language models*, 2023. arXiv: 2304.11619 [cs.CV]. [Online]. Available: https://arxiv.org/abs/2304.11619.

[24] T. Nguyen, J. Jewik, H. Bansal, P. Sharma, and A. Grover, *Climatelearn: Benchmarking machine learning for weather and climate modeling*, 2023. arXiv: 2307.01909 [cs.LG]. [Online]. Available: https://arxiv.org/abs/2307.01909.

[25] A. Srivastava, A. Rastogi, A. Rao, *et al.*, *Beyond the imitation game: Quantifying and extrapolating the capabilities of language models*, 2023. arXiv: 2206.04615 [cs.CL]. [Online]. Available: https://arxiv.org/abs/2206.04615.

[26] A. Talmor, J. Herzig, N. Lourie, and J. Berant, *Commonsenseqa: A question answering challenge targeting commonsense knowledge*, 2019. arXiv: 1811.00937 [cs.CL]. [Online]. Available: https://arxiv.org/abs/1811.00937.

[27] K. Sakaguchi, R. L. Bras, C. Bhagavatula, and Y. Choi, *Winogrande: An adversarial winograd schema challenge at scale*, 2019. arXiv: 1907.10641 [cs.CL]. [Online]. Available: https://arxiv.org/abs/1907.10641.

[28] B. Hawks, N. Tran, *et al.*, *Fast machine learning for science: Benchmarks and dataset*, 2022. [Online]. Available: https://arxiv.org/abs/2207.07958.

[29] B. Hawks, N. Tran, *et al.*, *Fast machine learning for science: Benchmarks and dataset*, 2022. [Online]. Available: https://arxiv.org/abs/2207.07958.

[30] B. Hawks, N. Tran, *et al.*, *Fast machine learning for science: Benchmarks and dataset*, 2022. [Online]. Available: https://arxiv.org/abs/2207.07958.

[31] Q. Wang *et al.*, *Boostr: A dataset for accelerator control systems*, 2021. [Online]. Available: https://arxiv.org/abs/2101.08359.

[32] J. Kvapil, G. Borca-Tasciuc, N. Tran, *et al.*, *Intelligent experiments through real-time ai: Fast data processing and autonomous detector control for sphenix and future eic detectors*, 2025. [Online]. Available: https://arxiv.org/abs/2501.04845.

[33] J. Weitz, D. Demler, L. McDermott, N. Tran, and J. Duarte, *Neural architecture codesign for fast physics applications*, 2025. [Online]. Available: https://arxiv.org/abs/2501.05515.

[34] B. Parpillon and N. Tran, *Smart pixels: In-pixel ai for on-sensor data filtering*, 2024. [Online]. Available: https://arxiv.org/abs/2406.14860.

[35] Y. Xiao and ERROR, *Braggnn: Fast x-ray bragg peak analysis using deep learning*, 2020. [Online]. Available: https://arxiv.org/abs/2008.08198.

[36] Anonymous, *4d-stem: Real-time ml for electron microscopy*, 2023. [Online]. Available: https://openreview.net/pdf?id=7yt3N0o0W9.

[37] J. Smith and J. Doe, *In-situ high-speed computer vision for plasma diagnostics*, 2023. [Online]. Available: https://arxiv.org/abs/2312.00128.

[38] W. Gao, J. Zhan, *et al.*, *Aibench: An industry standard internet service ai benchmark suite*, 2020. [Online]. Available: https://arxiv.org/abs/1908.08998.

[39] W. Gao, J. Zhan, *et al.*, *Bigdatabench: A scalable and unified big data and ai benchmark suite*, 2018. [Online]. Available: https://arxiv.org/abs/1802.08254.

[40] S. Farrell, M. Emani, *et al.*, *Mlperf hpc: A holistic benchmark suite for scientific machine learning on hpc systems*, 2021. [Online]. Available: https://arxiv.org/abs/2110.11466.

[41] M. S. W. Group, *Mlcommons science working group benchmarks*, 2023. [Online]. Available: https://github.com/mlcommons/science.

[42] E. Govorkova, E. Puljak, M. Pierini, *et al.*, "Lhc physics dataset for unsupervised new physics detection at 40 mhz," *Scientific Data*, 2022. DOI: 10.6084/m9.figshare.5046389. [Online]. Available: https://doi.org/10.5281/zenodo.5046389.

[43] C. Krause, B. Nachman, *et al.*, *Calochallenge 2022: A community challenge for fast calorimeter simulation*, 2024. [Online]. Available: https://arxiv.org/abs/2410.21611.

[44] P. W. Code, *Papers with code: Open machine learning benchmarks and leaderboards*, 2025. [Online]. Available: https://paperswithcode.com.

[45] Z. Xu, S. Escalera, *et al.*, "Codabench: Flexible, easy-to-use, and reproducible meta-benchmark platform," *Patterns*, vol. 3, no. 7, p. 100 543, 2022. DOI: 10.1016/j.patter.2022.100543.

[46] P. Luszczek *et al.*, "Sabath: Fair metadata technology for surrogate benchmarks," University of Tennessee, Tech. Rep., 2021.

[47] M. Takamoto, T. Praditia, *et al.*, "Pdebench: An extensive benchmark for scientific machine learning," in *NeurIPS Datasets and Benchmarks Track*, 2022. [Online]. Available: https://arxiv.org/abs/2210.07182.

[48] L. Zheng, L. Yin, *et al.*, *Sglang: Efficient execution of structured language model programs*, 2023. [Online]. Available: https://arxiv.org/abs/2312.07104.

[49] W. Kwon *et al.*, "Efficient memory management for large language model serving with pagedattention," in *SOSP 2023*, 2023.

[50] S. Mo, *Vllm performance dashboard*, 2024. [Online]. Available: https://simon-mo-workspace.observablehq.cloud/vllm-dashboard-v0/.

[51] K. G. Olivares, C. Challú, *et al.*, *Neuralforecast: User friendly state-of-the-art neural forecasting models*, PyCon US, 2022. [Online]. Available: https://github.com/Nixtla/neuralforecast.

[52] C. Challu, K. G. Olivares, *et al.*, "Nhits: Neural hierarchical interpolation for time series forecasting," in *AAAI 2023*, 2023.

[53] A. Garza, C. Challu, *et al.*, *Timegpt-1: A foundation model for time series*, 2023. [Online]. Available: https://arxiv.org/abs/2310.03589.

[54] E. G. Campolongo *et al.*, *Building machine learning challenges for anomaly detection in science*, 2025. [Online]. Available: https://arxiv.org/abs/2503.02112.

[55] E. G. Campolongo *et al.*, *Building machine learning challenges for anomaly detection in science*, 2025. [Online]. Available: https://arxiv.org/abs/2503.02112.

[56] E. G. Campolongo *et al.*, *Building machine learning challenges for anomaly detection in science*, 2025. [Online]. Available: https://arxiv.org/abs/2503.02112.

[57] G. Di Guglielmo, J. Campos, *et al.*, *End-to-end workflow for machine learning-based qubit readout with qick and hls4ml*, 2025. [Online]. Available: https://arxiv.org/abs/2501.14663.

[58] D. Rein, B. L. Hou, A. C. Stickland, *et al.*, *Gpqa: A graduate-level google-proof q and a benchmark*, 2023. [Online]. Available: https://arxiv.org/abs/2311.12022.

[59] K. X. Nguyen, F. Qiao, *et al.*, *Seafloorai: A large-scale vision-language dataset for seafloor geological survey*, 2024. [Online]. Available: https://arxiv.org/abs/2411.00172.

[60] R. Peterson, A. Tanelus, *et al.*, *Vocal call locator benchmark for localizing rodent vocalizations*, NeurIPS Poster, 2024. [Online]. Available: https://neurips.cc/virtual/2024/poster/97470.

[61] R. Bushuiev, A. Bushuiev, *et al.*, *Massspecgym: A benchmark for the discovery and identification of molecules*, NeurIPS Spotlight Poster, 2024. [Online]. Available: https://neurips.cc/virtual/2024/poster/97823.

[62] Y. Wang, T. Wang, *et al.*, *Urbandatalayer: A unified data pipeline for urban science*, NeurIPS Poster, 2024. [Online]. Available: https://neurips.cc/virtual/2024/poster/97837.

[63] W. Liu, R. Chen, *et al.*, *Delta squared-dft: Machine-learning corrected density functional theory for reaction energetics*, NeurIPS Poster, 2024. [Online]. Available: https://neurips.cc/virtual/2024/poster/97788.

[64] D. Patel, L. Zhao, *et al.*, *Large language models for crop science: Benchmarking domain reasoning and qa*, NeurIPS Poster, 2024. [Online]. Available: https://neurips.cc/virtual/2024/poster/97570.

[65] X. Zhong, Y. Gao, *et al.*, *Spiqa-llm: Evaluating llm adapters on scientific figure qa*, NeurIPS Poster, 2024. [Online]. Available: https://neurips.cc/virtual/2024/poster/97575.